

Instruction Set Architecture (ISA - Ch 2)

Dr. Rajib Ranjan Maiti (Mtech and PhD at IIT Kharagpur)
CSIS, BITS-Pilani, Hyderabad

Supporting Procedures in Computer Hardware

- A procedure follow these six steps:
 - 1. Put parameters in a place where the procedure can access them.
 - 2. Transfer control to the procedure.
 - 3. Acquire the storage resources needed for the procedure.
 - 4. Perform the desired task.
 - 5. Put the result value in a place where the calling program can access it.
 - 6. Return control to the point of origin, since a procedure can be called from several points in a program.

- MIPS convention for procedure call :
 - \$a0–\$a3: four registers to pass arguments
 - \$v0–\$v1: two registers to return values
 - \$ra: one register to hold return address

Supporting Procedures in Computer Hardware

- Lets Consider C procedure:
 - `int leaf_example (int g, int h, int i, int j) {`
 - `int f;`
 - `f = (g + h) - (i + j);`
 - `return f;`
 - `}`
- What is the compiled MIPS assembly code?

Supporting Procedures in Computer Hardware

- Lets Consider C procedure:
 - `int leaf_example (int g, int h, int i, int j) {`
 - `int f;`
 - `f = (g + h) - (i + j);`
 - `return f;`
 - `}`
- What is the compiled MIPS assembly code?

- The arguments : g, h, i, and j are in \$a0, \$a1, \$a2, and \$a3,
- Local variable : f is in \$s0
- MIPS Code:

- `leaf_example: addi $sp, $sp, -12 #adjust stack`
 - `sw $t1, 8($sp) # save $t1 for use afterwards`
 - `sw $t0, 4($sp) # save $t0 for use afterwards`
 - `sw $s0, 0($sp) # save $s0 for use afterwards`
- `add $t0,$a0,$a1 # register $t0 contains g + h`
- `add $t1,$a2,$a3 # register $t1 contains i + j`
- `sub $s0,$t0,$t1 # f = $t0 - $t1, which is (g + h) - (i + j)`
- `add $v0,$s0,$zero # returns f ($v0 = $s0 + 0)`
- `lw $s0, 0($sp) # restore $s0 for caller`
- `lw $t0, 4($sp) # restore $t0 for caller`
- `lw $t1, 8($sp) # restore $t1 for caller`
- `addi $sp,$sp,12 # adjust stack to delete 3 items`
- `jr $ra # jump back to calling routine`

Communicating with People

- **ASCII versus Binary Numbers**
- How much does storage increase if the number 1 billion is represented in ASCII versus a 32-bit integer?
 - Ans. One billion is 1,000,000,000, so it would take 10 ASCII digits, each 8 bits long.
- How MIPS deals with byte?
 - **lb \$t0, 0(\$sp)** # Read byte from memory
 - **sb \$t0, 0(\$gp)** # Write byte to memory
- Consider following C codes
 - ```
void strcpy (char x[], char y[]) {
 int i;
 i = 0;
 while ((x[i] = y[i]) != '\0') /* copy &
 test byte */
 i += 1;
}
```
- What is the MIPS assembly code?

# MIPS Addressing for I-Type instructions

- Although constants can fit in a 16-bit field,
  - sometimes they are bigger.
- Consider
  - **lui \$t0, 255**      **# load upper immediate - \$t0 is register 8:**
- 32 bit format of the instruction

| op     | rs              | rd     | Immid.              |
|--------|-----------------|--------|---------------------|
| 001111 | 00000           | 01000  | 0000 0000 1111 1111 |
| 6 bits | 5 bits (unused) | 5 bits | 16 bits             |

- Content of \$t0 after execution

|                     |                     |
|---------------------|---------------------|
| 0000 0000 1111 1111 | 0000 0000 0000 0000 |
| Upper 16 bits       | Lower 16 bits       |

# MIPS Addressing I-Type Instructions

- bne \$s0,\$s1,Exit # go to Exit if \$s0  $\neq$  \$s1

| Op     | rs     | rt     | Immid.  |
|--------|--------|--------|---------|
| 5      | 16     | 17     | Exit    |
| 6 bits | 5 bits | 5 bits | 16 bits |

# MIPS Addressing J-Type Instructions

- **Addressing in Branches and Jumps**

- j 10000                      # go to location 10000

| Op     | Immid.  |
|--------|---------|
| 2      | 10000   |
| 6 bits | 26 bits |



# Calculating Branch target address

- MIPS assembler code:
  - #variable i is in \$s3 and base is in \$s6
  - Loop: sll \$t1,\$s3,2           # Temp reg \$t1 = 4 \* i
  - add \$t1,\$t1,\$s6       # \$t1 = address of save[i]
  - lw \$t0,0(\$t1)       # Temp reg \$t0 = save[i]
  - bne \$t0,\$s5, Exit   # go to Exit if save[i] ≠ k
  - addi \$s3,\$s3,1       # i = i + 1
  - j Loop               # go to Loop
  - Exit:
- If we assume we place the loop starting at location 80000 in memory, what is the MIPS machine code for this loop?

|       |     |                                                          |    |                                     |   |    |
|-------|-----|----------------------------------------------------------|----|-------------------------------------|---|----|
| 80000 | 0   | 0                                                        | 19 | 9                                   | 2 | 0  |
| 80004 | 0   | 9                                                        | 22 | 9                                   | 0 | 32 |
| 80008 | 35  | 9                                                        | 8  | 0                                   |   |    |
| 80012 | 5   | 8                                                        | 21 | 2 #Exit is +2 away from next instr. |   |    |
| 80016 | 8   | 19                                                       | 19 | 1                                   |   |    |
| 80020 | 2   | 20000 #Loop 800000, so, field value = 800000 / 4 = 20000 |    |                                     |   |    |
| 80024 | ... |                                                          |    |                                     |   |    |

| Mnemonic | Meaning                                | Type | Opcode | Funct |
|----------|----------------------------------------|------|--------|-------|
| add      | Add                                    | R    | 0x00   | 0x20  |
| addu     | Add Unsigned                           | R    | 0x00   | 0x21  |
| and      | Bitwise AND                            | R    | 0x00   | 0x24  |
| div      | Divide                                 | R    | 0x00   | 0x1A  |
| divu     | Unsigned Divide                        | R    | 0x00   | 0x1B  |
| jalr     | Jump and Link Register                 | R    | 0x00   | 0x09  |
| jr       | Jump to Address in Register            | R    | 0x00   | 0x08  |
| mfhi     | Move from HI Register                  | R    | 0x00   | 0x10  |
| mthi     | Move to HI Register                    | R    | 0x00   | 0x11  |
| mflo     | Move from LO Register                  | R    | 0x00   | 0x12  |
| mtlo     | Move to LO Register                    | R    | 0x00   | 0x13  |
| mfco     | Move from Coprocessor 0                | R    | 0x10   | NA    |
| mult     | Multiply                               | R    | 0x00   | 0x18  |
| multu    | Unsigned Multiply                      | R    | 0x00   | 0x19  |
| nor      | Bitwise NOR (NOT-OR)                   | R    | 0x00   | 0x27  |
| xor      | Bitwise XOR (Exclusive-OR)             | R    | 0x00   | 0x26  |
| or       | Bitwise OR                             | R    | 0x00   | 0x25  |
| slt      | Set to 1 if Less Than                  | R    | 0x00   | 0x2A  |
| sltu     | Set to 1 if Less Than Unsigned         | R    | 0x00   | 0x2B  |
| sll      | Logical Shift Left                     | R    | 0x00   | 0x00  |
| srl      | Logical Shift Right (0-extended)       | R    | 0x00   | 0x02  |
| sra      | Arithmetic Shift Right (sign-extended) | R    | 0x00   | 0x03  |
| sub      | Subtract                               | R    | 0x00   | 0x22  |
| subu     | Unsigned Subtract                      | R    | 0x00   | 0x23  |

| Mnemonic | Meaning                                  | Type | Opcode | Funct |
|----------|------------------------------------------|------|--------|-------|
| j        | Jump to Address                          | J    | 0x02   | NA    |
| jal      | Jump and Link                            | J    | 0x03   | NA    |
| addi     | Add Immediate                            | I    | 0x08   | NA    |
| addiu    | Add Unsigned Immediate                   | I    | 0x09   | NA    |
| andi     | Bitwise AND Immediate                    | I    | 0x0C   | NA    |
| beq      | Branch if Equal                          | I    | 0x04   | NA    |
| blez     | Branch if Less Than or Equal to Zero     | I    | 0x06   | NA    |
| bne      | Branch if Not Equal                      | I    | 0x05   | NA    |
| bgtz     | Branch on Greater Than Zero              | I    | 0x07   | NA    |
| lb       | Load Byte                                | I    | 0x20   | NA    |
| lbu      | Load Byte Unsigned                       | I    | 0x24   | NA    |
| lhu      | Load Halfword Unsigned                   | I    | 0x25   | NA    |
| lui      | Load Upper Immediate                     | I    | 0x0F   | NA    |
| lw       | Load Word                                | I    | 0x23   | NA    |
| ori      | Bitwise OR Immediate                     | I    | 0x0D   | NA    |
| sb       | Store Byte                               | I    | 0x28   | NA    |
| sh       | Store Halfword                           | I    | 0x29   | NA    |
| slti     | Set to 1 if Less Than Immediate          | I    | 0x0A   | NA    |
| sltiu    | Set to 1 if Less Than Unsigned Immediate | I    | 0x0B   | NA    |
| sw       | Store Word                               | I    | 0x2B   | NA    |

# MIPS Pseudo-Instruction

abs  
blt  
bgt  
ble  
neg  
negu  
not  
bge  
li  
la  
move  
sge  
sgt

| Pseudo-Instruction  | MIPS Instructions                                                                          |
|---------------------|--------------------------------------------------------------------------------------------|
| abs \$1, \$2        | addu \$1, \$2, \$0<br>bgez \$2, 8 (offset=8 → skip 'sub' instruction)<br>sub \$1, \$0, \$2 |
| blt \$8, \$9, label | slt \$1, \$8, \$9<br>bne \$1, \$0, label                                                   |
| li \$8, 0x3BF20     | lui \$at, 0x0003<br>ori \$8, \$at, 0xBF20                                                  |
| la \$a0, address    | lui \$at, 4097 (0x1001 → upper 16 bits of \$at).<br>ori \$a0, \$at, disp                   |
| move \$1, \$2       | add \$1, \$2, \$0                                                                          |
| sge \$1, \$8, \$9   | addiu \$9, \$9, -0x01<br>slt \$1, \$9, \$8                                                 |

# MIPS addressing modes

- Addressing mode
  - Multiple forms of addressing

## The MIPS addressing modes

**Immediate addressing:**  
the operand  
is a constant  
**within the  
instruction  
itself**

### 1. Immediate addressing



# MIPS addressing modes

- Addressing mode
  - Multiple forms of addressing

## The MIPS addressing modes

**Immediate addressing:**  
the operand is a constant **within the instruction itself**

**Register addressing:**  
the operand is a **register**

### 2. Register addressing



Registers

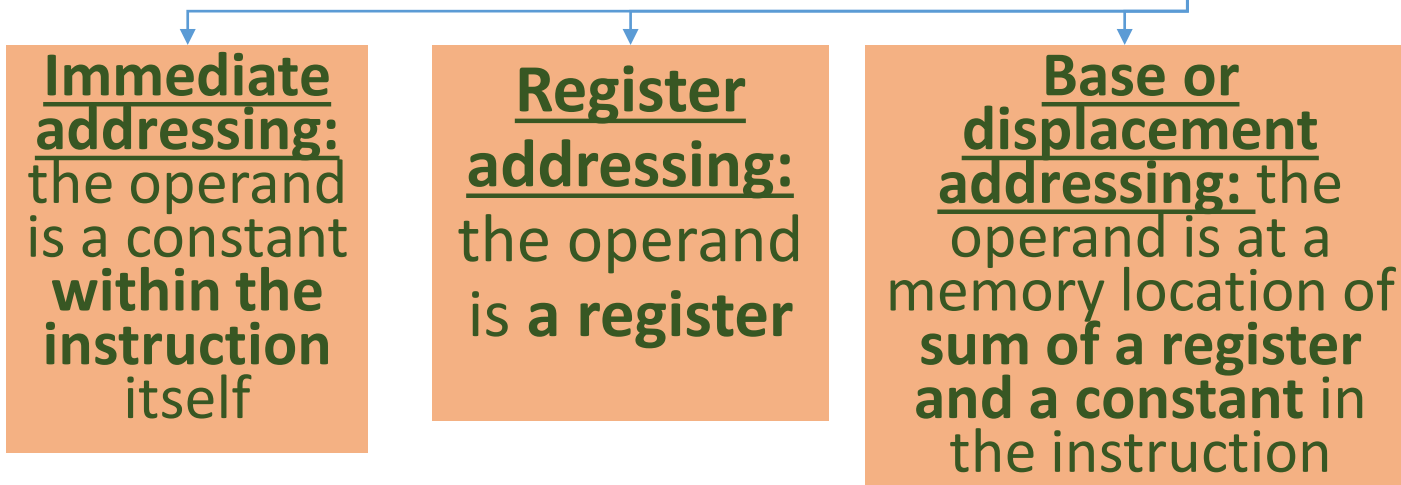
Register




# MIPS addressing modes

- Addressing mode
  - Multiple forms of addressing

## The MIPS addressing modes



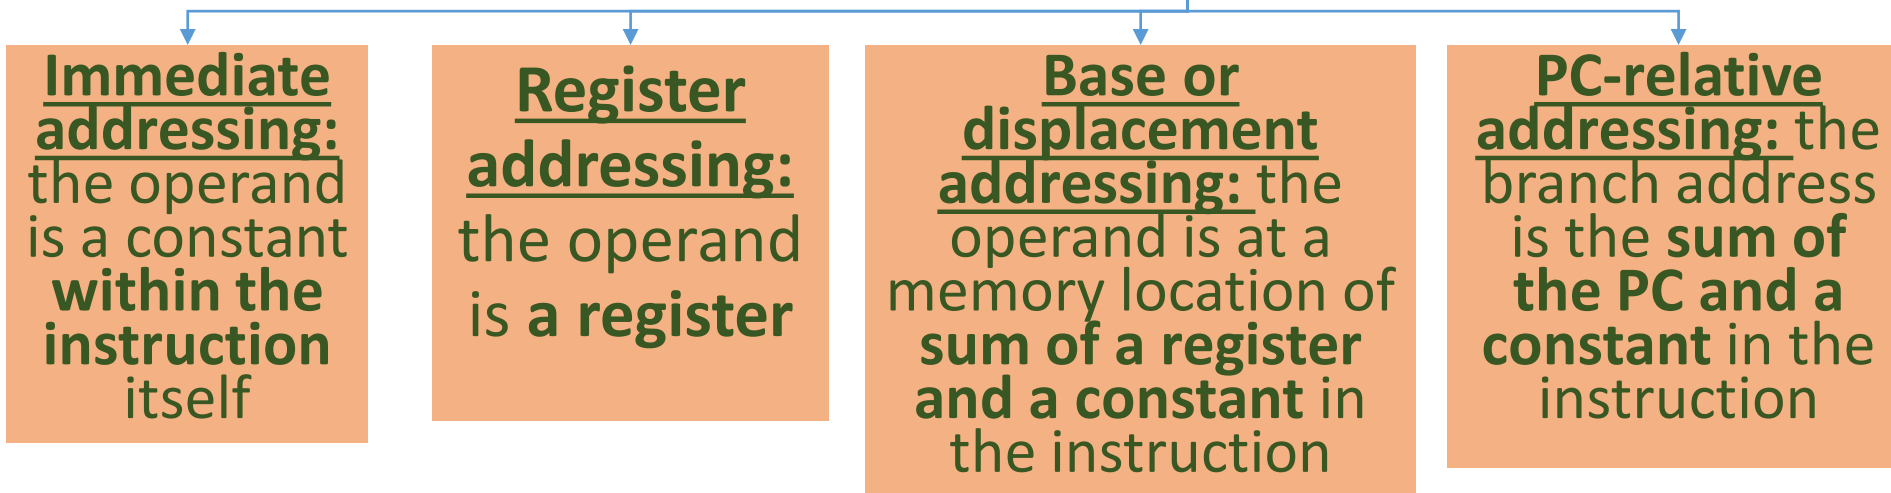
### 3. Base addressing




# MIPS addressing modes

- Addressing mode
  - Multiple forms of addressing

## The MIPS addressing modes



4. PC-relative addressing



# MIPS addressing modes

- Addressing mode
  - Multiple forms of addressing

## The MIPS addressing modes

**Immediate addressing:** the operand is a constant **within the instruction itself**


**Register addressing:** the operand is a **register**

**Base or displacement addressing:** the operand is at a memory location of **sum of a register and a constant** in the instruction

**PC-relative addressing:** the branch address is the **sum of the PC and a constant** in the instruction

**Pseudodirect addressing:** the jump address is the **26 bits of the instruction** concatenated with the upper bits of the PC

5. Pseudodirect addressing





# Parallelism and Instructions: Synchronization

- *Atomic exchange or atomic swap*
  - A value in a register  $\leftrightarrow$  A value in memory
- Build a simple **lock** where
  - Lock = 0  $\rightarrow$  the lock is free
  - Lock = 1  $\rightarrow$  the lock is unavailable.
- The value returned from the exchange instruction is
  - 1 if some other processor had already claimed access
  - 0 otherwise.

# Parallelism and Instructions: Synchronization

- In MIPS, this pair of instructions includes
  - a special load instruction
    - *load linked (ll)*
  - a special store
    - *store conditional (sc)*
- These instructions are used in sequence:
  - if the contents of the memory location specified by the *ll* are changed before the *sc* to the same address occurs, then *sc* fails

# Concurrency Problem

- What will be the possible values in `*$s0` when the following code executed concurrently by two threads?
  - `# *($s0) = 100` -- initial value
  - `lw $t0, 0($s0)`
  - `addi $t0, $t0, 1`
  - `sw $t0, 0($s0)`
- Ans:
  - **101 or 102**
  - 100, 101 or 102
  - 100 or 101
  - 102

# Concurrency Problem

- A possible solution
  - Lock (a.k.a. busy wait)
    - Get\_lock: # \$s0 -> addr of lock
    - addiu \$t1, \$zero, 1 #t1 = Locked value
    - Loop:   lw \$t0, 0(\$s0) #load lock
    - bne \$t0, \$zero, Loop # loop if locked
    - Lock:   sw \$t1, 0(\$s0) # Unlocked, so lock
  - Unlock
    - Unlock: sw \$zero, 0(\$s0)

Any problems with this?

An  
execution  
scenario

## ❑ Thread 1

```
 addiu $t1,$zero,1
Loop: lw $t0,0($s0)

 bne $t0,$zero,Loop

Lock: sw $t1,0($s0)
```

## ❑ Thread 2

```
 addiu $t1,$zero,1
Loop: lw $t0,0($s0)

 bne $t0,$zero,Loop

Lock: sw $t1,0($s0)
```

Time

Both threads think they have set the lock!  
Exclusive access not guaranteed!

# A solution using ll and sc (atomic swap)

- *ll* returns the initial value, *sc* returns 1 only if it succeeds
- For an atomic exchange between  $0(\$s1)$ , i.e.,  $M[\$s1 + 0]$ , and  $\$s4$ :
  - **AGAIN:** `addi $t0, $zero, $s4`       $\#\$t0 \leftarrow \$s4$
  - `ll $t1, 0($s1)`               $\#\$t1 \leftarrow 0(\$s1)$  using *ll*
  - `sc $t0, 0($s1)`               $\#0(\$s1) \leftarrow \$t0$  using *sc*
  - `beq $t0, $zero, AGAIN`       $\#$ try again if *sc* fails
  - `add $s4, $zero, $t1`         $\#\$s4 \leftarrow \$t1$
- Any time a processor intervenes and modifies the value in memory between the *ll* and *sc* instructions,
  - *sc* returns 0 in  $\$t0 \rightarrow$  try again.
- At the end of this sequence,
  - the contents atomically exchanged between  $\$s4$  and  $0(\$s1)$

# Example of Test-and-Set using ll and sc

- In a single atomic operation:
  - **Test** : to see if a memory location is set (contains a 1)
  - **Set** : the memory location to 1, if the location has not contained 1 when tested, i.e., contained a zero
  - **Otherwise** : the **Set** failed, so the program can try again
  - While accessing, no other instruction can modify the memory location, including other Test-and-Set instructions
- Useful for implementing lock operation

# Example of Test-and-Set using ll and sc

- MIPS sequence for implementing a T&S at (\$s1)
  - Try: addiu \$t0, \$zero, 1                      # \$t0  $\leftarrow$  1
  - ll \$t1, 0(\$s1)                              # \$t1  $\leftarrow$  M[\$s1 + 0]
  - bne \$t1, \$zero, Try                        # if \$t1  $\neq$  0 then Try again
  - sc \$t0, 0(\$s1)                              # Otherwise, M[\$s1+0]  $\leftarrow$  \$t0
  - beq \$t0, \$zero, Try                        # if \$t0 == 0 then Try again
  - Locked:
    - critical section
    - sw \$zero, 0(\$s1)                          # release lock



# A C Sort Example to Put It All Together

# A C Sort Example to Put It All Together

- The procedure

- `void swap(int v[], int k) {`
  - `int temp;`
  - `temp = v[k];`
  - `v[k] = v[k+1];`
  - `v[k+1] = temp;`
- `}`

- To translate from C to assembly language by hand, follow the general steps of:

- 1. Allocate registers to program variables.
- 2. Produce code for the body of the procedure.
- 3. Preserve registers across the procedure invocation.

# Steps to get MIPS code

- The variables:

- The arguments:

- $v$  in  $\$a0$
    - $k$  in  $\$a1$

- The local variable:

- $temp$  in  $\$t0$

- Body of the procedure

- `sll $t1, $a1, 2` # reg  $\$t1 = k * 4$
  - `add $t1, $a0, $t1` # reg  $\$t1 = v + (k * 4)$
  - # reg  $\$t1$  has the address of  $v[k]$
  - 
  - `lw $t0, 0($t1)` # reg  $\$t0$  (temp) =  $v[k]$
  - `lw $t2, 4($t1)` # reg  $\$t2 = v[k + 1]$
  - # refers to next element of  $v$
  - 
  - `sw $t2, 0($t1)` #  $v[k] = \text{reg } \$t2$
  - `sw $t0, 4($t1)` #  $v[k+1] = \text{reg } \$t0$  (temp)

- Return of caller

- `jr $ra` # return to calling routine

# Arrays versus Pointers

- `clear1(int array[], int size) {`
  - `int i;`
  - `for (i = 0; i < size; i += 1)`
    - `array[i] = 0;`
- `}`
- Assumptions:
  - Base Address of Array in `$a0`
  - Size of array in `$a1`
  - `i` in `$t0`.

- Initialization
  - `move $t0,$zero`                      `# i = 0 (register $t0 ← 0)`
- To set `array[i]` to 0
  - `loop1: sll $t1,$t0,2`                      `# $t1 = i * 4`
  - `add $t2,$a0,$t1`                      `# $t2 = address of array[i]`
  - `sw $zero, 0($t2)`                      `# array[i] = 0`
  - `addi $t0,$t0,1`                      `# i = i + 1`
  - `slt $t3,$t0,$a1`                      `# $t3 = (i < size)`
  - `bne $t3,$zero,loop1`                      `# if $t3 !=0 then go to loop1`

# Arrays versus Pointers: Pointer Version

- `clear2(int *array, int size) {`
  - `int *p;`
  - `for (p = &array[0]; p < &array[size]; p = p + 1)`
    - `*p = 0;`
- `}`

- Assumptions:

- Base Address of Array in `$a0`
- Size of array in `$a1`
- `i` in `$t0`.

- Initialization:

- `move $t0,$a0` # `p = address of array[0]`

- To set `*p` to 0

- `loop2: sw $zero,0($t0)` # `Memory[p] = 0`
- `addi $t0,$t0,4` # `p = p + 4`
- `sll $t1,$a1,2` # `$t1 = size * 4`
- `add $t2,$a0,$t1` # `$t2 = address of array[size]`
- `slt $t3,$t0,$t2` # `$t3 = (p < &array[size])`
- `bne $t3,$zero,loop2` # if `(p < &array[size])` go to `loop2`

# Fallacies and Pitfalls

# Fallacies: commonly held misconceptions

- *More powerful instructions mean higher performance.*
  - Find a counterexample or counter argument
- Prefix in Intel x86
  - Part of the power of the Intel x86 is the prefixes that can modify the execution of the following instruction.
  - One prefix can repeat the following instruction until a counter counts down to 0.
  - Thus, to move data in memory, it would seem that the natural instruction sequence is to use move with the repeat prefix to perform 32-bit memory-to-memory moves.
- An alternative method,
  - Use the standard instructions found in all computers,
    - First load the data into the registers
    - Then store the registers back to memory.
  - This second version of this program, with the code replicated to reduce loop overhead, copies at about 1.5 times as fast.
- Another alternative method
  - Use the larger floating-point registers instead of the integer registers of the x86, copies at about 2.0 times as fast than the complex move instruction.

# Fallacies: commonly held misconceptions

- *Write in assembly language to obtain the highest performance.*
  - Find a counterexample
- When compilers were poor at register allocation, so programmers use certain hints in high level programming language, like C, to achieve a better performance.
- Today's C compilers generally ignore such hints, because the compiler does a better job at allocation than the programmer does.
- One of the few widely accepted axioms of software engineering is that coding takes longer if you write more lines, and it clearly takes many more lines to write a program in assembly language.
- Moreover, once it is coded, the next danger is that it will become a popular program. Such programs always live longer than expected, meaning that someone will have to update the code over several years and make it work with new releases of operating systems and new models of machines.



# Fallacies: commonly held misconceptions

- *The importance of commercial binary compatibility means successful instruction sets don't change.*
  - Find a counterexample

## x86 ISA over time




FIGURE 2.43 Growth of x86 instruction set over time. While there is clear technical value to some of these extensions, this rapid change also increases the difficulty for other companies to try to build compatible processors.

# Pitfalls: easily made mistakes

- *Forgetting that sequential word addresses in machines with byte addressing do not differ by one.*
  - Because of generalizations of principles that are only true in a limited context
- Many an assembly language programmer has toiled over errors made by assuming that the address of the next word can be found by incrementing the address in a register by one instead of by the word size in bytes.

# Pitfalls: easily made mistakes

- *Using a pointer to an automatic variable outside its defining procedure.*
  - Because of generalizations of principles that are only true in a limited context
- pass a result from a procedure that includes a pointer to an array that is local to that procedure
- the memory that contains the local array will be reused as soon as the procedure returns
- Pointers to automatic variables can lead to chaos.

*Thank You*